CS111 Chapter 2.1 Outline #3

Overview

- Three main, fundamental abstractions
 - Memory (number of data items to remember)
 - Interpreter (number of steps the interpreter must execute)
 - Communication Link (number of messages to communicate)
- These are used to organize physical hardware structures
 - They supply functions of recall, processing, and communication
 - Widely useful with understandably-simple interface semantics
- System designers rearrange/repackage these abstractions, layering them into more customized ones
- Reference primary method by which abstract components in the system interact
 - Usual method for components to connect is by "name," which appear in their interfaces
 - Memory stores/retrieves by name; interpreter manipulates by name; names identify communication links

Memory (storage)

- System component that remembers data values for computation
- o Simple abstract model for all memory devices
 - WRITE(name, value)
 - Value to be remembered; it can be recalled by its name
 - value ←READ(name)
 - Memory device remembers value associated with given name
- Volatile (consumes energy to retain energy; interrupted power means information loss)
 - When connected to a battery or uninterruptable power supply, it becomes durable (able to remember for some period, or "durability")
- Non-volatile or stable storage (always retains content; power means READ operations can be run)
- Both types of memory face decay, so every such device has a durability
- o Memory devices don't usually name, READ, or WRITE values of arbitrary size
 - Operates contiguous arrays of (fixed) bits, or bytes (higher-layer uses record, segments, or files)
- o Memory/storage cell: unit of physical layer memory written or read
- Read/Write Coherence & Before-Or-After Atomicity
 - Read/Write Coherence: result of READing a named cell is the same as the most recent WRITE to that cell
 - o Before-Or-After Atomicity: every READ/WRITE result occurs distinctly before/after other R/W operations
- Threats against Read/Write Coherence and Atomicity
 - Concurrency: concurrent READ/WRITEs on same cell needs arbitration so one occurs completely before
 - o Remote storage: physical distance and lag may make unclear which operation was most recent
 - Performance enhancements: processors might reorder operations for optimization, thus destroying concurrency; programmers can force code to run concurrently (SYNCHRONIZED; memory barrier)
 - Cell size incommensurate with value size: large values may take up multiple cells, and a concurrent READ/WRITE may cause "write tearing," or only some cells of that value are updated; small values may share a single cell, and concurrent operations might overwrite the other; atomic coordination is needed
 - Tradeoff between efficiency and R/W Coherence and Atomicity in a system design
- Memory latency (access time)
 - o Time taken for a READ/WRITE to complete
 - o Magnetic disk memory latency depends on the device's mechanical state at request instance
 - If the disk rotates past a second nearby sector that the user wants to access, there may be a thousand-time delay as we wait for the next rotation under the read head
 - Max transfer rate to a disk is much larger than that of random sectors
 - Random access memory (RAM)
 - Latency of randomly-chosen memory cells is about the same as those chosen in the best order
 - Behavior unlike other mechanical memory, like optical disks or magnetic tapes/disks
 - For non-random memory, it is more worthwhile to READ/WRITE huge blocks
- Memory names/addresses
 - o Memory cells are generally named by the geometric coordinates of their physical storage location

- Geometric coordinates are sometimes mapped in consecutive integers called "addresses," which form the address space of the memory device
- Such a memory system is called location-addressed memory
- Associative memory: memory system that accepts unconstrained names (unlike consecutive integers)
- Cache: device that remembers the result of an expensive computation in order to reuse it
 - Often associative hardware/software
- Exploiting memory abstraction (RAID)
 - Redundant Array of Independent/Inexpensive Disks
 - An example of storage abstraction and modularity with two goals
 - Improved performance by reading/writing disks concurrently
 - Improved durability by writing information on more than one disk
 - Various configurations, such as allowing concurrent read/write, error correcting, or copying data

Interpreter

- Active elements in a computer system for performing computational actions
- o Abstraction of three components
 - Instruction reference: tells interpreter where to find next instruction
 - Repertoire: set of ready-to-perform actions by interpreter when instruction is found at location
 - Environment reference: tells interpreter where to find environment (or state on which to act)
- o Interpreter normally sequences through the program, finds environment by environment reference, takes the program instruction from that, then performs the action by changing some data; repeat
- Interrupts (handler) might provide the interpreter with code rather than the program's instruction code
- Multiple interpreters are often asynchronous (run on uncoordinated clocks) and run at different rates
 - Contributes to coherence/atomicity problems in memory

Processor

- Implementation of an interpreter; instruction reference is a program counter in the processor's fast memory processor; may be wired directly to memory; addresses in the program counter are names in that memory's address space
- The repertoire includes instructions that express computations or move data
- A stack is provided for implementing procedure calls
 - Caller pushes callee (arguments) on the stack; when the callee returns, stack pops back
 - Stack pointer register holds memory address of top of stack for efficiency
- Interrupts may occur due to an instruction that cannot be implemented or an external signal that needs attention
 - Exception handler may take over in the former case
 - In the latter case, interrupt handler may do some work, then return control

• Interpreter Layers

- o Interpreters are organized in layers (a full application might have 4 or 5 layers)
- Lowest layer is hardware with primitive repertoire; successive layers have more functionality
- Example: hardware \rightarrow interpreter \rightarrow program \rightarrow (user)
- o Goal: ensure designers that the layer under another layer either completes instruction or does nothing
- Assume interpreters are atomic

• Communication Link

- Provides a means for information to move between physically-separated components
- SEND(link name, outgoing message buffer)
 - An array of bits (a message) is transmitted across link name
- RECEIVE(link_name, incoming_message_buffer)
 - Some buffer is specified to hold the incoming message
 - Higher interfaces can RECEIVE from lower ones (lower ones can "upcall," or deliver)
- The link_name argument is a possible communication link; some links are multiply-attached networks of links with another recipient-choosing method
- Differences from READ/WRITE make the semantics and implementation rather different
 - o Parameters make SEND/RECEIVE times unpredictable; if asynchronous, time is unknown in advance
 - Environment threatens integrity of data transfer; in some cases, the message might fail to deliver