

CS111 Chapter 4-5 Outline #4

Chapter 4, Section 4.1-4.1.2 (pages 147-162), Chapter 5, Section 5.1 (pages 199-210), Section 5.3-5.3.4 (pages 230-237)

- Enforced Modularity
 - Idea: errors can propagate even through modules
 - Client/service organization comprises modules that communicate by message only
 - Messages are the only way...
 - To provide a service, preventing the programmer from convention violations
 - For errors to propagate, thus allowing independent failure
 - For an attacker to penetrate module; checking messages can block attacks
 - Effective implementation: run client/service module on separate computers with a communication wire
 - Prevents errors from propagating
- Client/Service Organization
 - Soft modularity limits interactions of certain modules to their interfaces, but implementation errors might bypass the interfaces (example: caller and callee use the same address space and stack)
 - Need hard boundaries between modules (client/server: sweeping simplification regarding interfaces)
 - Errors propagate only via messages
 - Errors can be found via message checking
 - MEASURE and GET_TIME example of soft modularity
 - In MEASURE, caller and callee don't have a full contract that addresses errors
 - They modify shared arguments and their own variables in stack; callee could potentially corrupt caller's area of the stack (since the caller leaves the stack pointer where it is)
 - Callee returns to where caller specifies (even if it makes a mistake)
 - Callee stores values in register R0; if it's misplaced, the caller will grab a wrong value
 - If callee divides by zero, the caller might terminate (known as "fate sharing")
 - Caller and callee are only allowed to modify global variables shared between them (otherwise, other modules might fail in computation or crash)
 - Soft modularity is found in specifications, but it is not enforced
 - Enforced modularity by an external mechanism is desired; it limits the user of interactions
 - Benefitted with a high-level language due to compiler/runtime system (Java, C#)
 - Other shortcomings that demonstrate the importance of client/service
 - Different programming languages in modules
 - Implementation error in interpreter can lead to incorrectly-restored registers, etc.
- Client/Service Components
 - Client: module that initiates request (builds "request" message containing data to send to service)
 - Service: modules that responds (takes arguments from message, executes operations, sends back "response/reply" message, waits for next request)
 - Designers use "message timing diagrams" to represent interactions
 - Two separate computers (can be geographically remote)
 - Can avoid absolute failure during power outages; costs more
 - A service of multiple computers is called a server (more fault tolerance)
 - Marshaling: the conversion of arguments between the client and service computer(s) (like endian-ness)
 - Unmarshaling converts it back to a language object
 - Client/service organization separates functions (abstraction) and enforces it (enforced modularity)
 - Benefits:
 - No shared state other than messages
 - Transaction is at arm's length; arguments can be marshaled and thus checked
 - Client protects itself from even failing services by ignoring infinite loops after a set time
 - Encourages explicit, well-defined interfaces (messages that are specific)
 - If the service has a disaster, the client has a controlled problem (only problem is lack of result)
 - Sweeping simplification: only messages are passed (and essentially between a firewall)
 - Detriments:
 - Service might still return an erroneous result

- System designer and client need a good interface for catching certain problems
- To use or not to use?
 - Tradeoff between ease of data access and ease of error propagation
 - Client and service units become separate programs that are harder to manipulate shared data
 - Determine plan for recovery such that fate sharing is reduced
 - Example: the world wide web (browser is client, site is service)
- Enforcing Modularity with Virtualization
 - Client/service: enforces modularity to prevent errors, achieves security, better fault tolerance
 - Unfortunately, uses a computer for each module
 - Virtualization with three new abstractions
 - SEND and RECEIVE using bounded buffers
 - Virtual memory
 - Threads
 - Client/Service Organization via Virtualization
 - Use one physical computer to run multiple virtual computers, each running a module
 - Virtualization
 - Multiplexing: program simulates interface by multiplexing a physical instance into many virtual objects
 - Example: server into web sites, processors into threads
 - Aggregating: program simulates interface by aggregating many physical instances into a large virtual object
 - Example: disks into RAID, communication channels into bonding
 - Emulation: implement virtual object from a different kind of physical object
 - Example: disk to RAM, Mac to virtual PC
 - Virtual computers don't enforce modularity as well (less fault tolerance, i.e. to power failures)
 - Threads (virtualized processors)
 - Abstraction that encapsulates execution state of an active computation
 - Reference to next program step
 - References to environment
 - Threads can be stopped and restarted later; used for virtualizing processors
 - One thread (serial executions) follows the principle of least astonishment
 - More than one thread (say, per device) allows concurrent operations
 - Thread abstraction implemented by thread manager
 - Multiplexing multiple threads on limited physical processors to prevent interference
 - Threads allow interrupts to be processed concurrently
 - Exceptions are interrupts on current thread
 - Virtual memory
 - Sharing memory has benefits, but can be easy to make mistakes
 - Threads of one module cannot overwrite another module's data
 - Virtual memory gives each module its own virtual address space and virtual addresses
 - These are controlled by virtual memory manager
 - SEND/RECEIVE with bounded buffers (virtualized communication links)
 - Calling SEND fills buffer until it is full; then it waits until the bounded buffer has more space
 - Calling RECEIVE waits (the calling thread waits) if the buffer is empty
 - OS Interface
 - Memory
 - Create/delete address space
 - Allocate/free block
 - Interpreter
 - Allocate/delete/exit thread
 - Yield, await, advance, etc
 - Communication
 - Allocate/deallocate bounded buffer

- Send/receive
- Emulation and Virtual Machines
 - Emulation simulates physical hardware so that the emulated hardware can run the same software the physical hardware can run
 - Microcode interpreter inside a processor can simulate instructions of other processors/instructions from other processors
 - Allows software development before chip is manufactured
 - Emulation is slow because interpreting emulated machine instructions has substantial overhead
 - Costs up to a factor of 10 in performance
 - Fast emulation uses virtual machines; physical processor implements many virtual instances
 - Less portability
 - Better performance
- Enforcing Modularity in Memory
 - Bounded buffers take advantage of threads sharing same physical memory
 - Enforced modularity means threads of one module cannot overwrite another module's data
 - Enforcing Modularity with Domains
 - Restrict thread references to a domain, or contiguous range of memory addresses
 - Each processor has a domain register to restrict memory references to that domain
 - Memory manager (a special interpreter) checks that each memory reference is between the range within the domain register
 - Otherwise, triggers a memory reference exception
 - Shortcomings:
 - Threads might need more than one domain
 - Threads should not be able to change its own domain
- Controlled Sharing Using Several Domains
 - Allow threads to have several domains; processors have several domain registers
 - Multiple threads can use shared bounded buffer domain, but not other (private domain) references
 - Interface:
 - $\text{base_address} \leftarrow \text{ALLOCATE_DOMAIN}(\text{size})$: allocate new domain, return base address
 - $\text{MAP_DOMAIN}(\text{base_address})$: add domain starting at base_address to thread's domains
 - Permissions can allow more controlled sharing
 - $\text{MAP_DOMAIN}(\text{base_address}, \text{permission})$
 - Load bounds from domain table into calling thread's registers with specified permission
 - LOAD requires READ permission; STORE requires WRITE permissions
 - Must be within domain with permissions
 - For domain register's the manager calls CHECK_DOMAIN (with parameters: address request, permissions needed, domain register) to determine whether to issue the request
 - Demultiplexing memory reference exceptions can be hardware or software
 - Hardware: illegal memory reference/permission error exception
 - Software: memory manager signals memory reference exception
 - Only some permission combinations are possible: R, R/W, R/E, R/W/E
- More Enforced Modularity with Kernel and User Mode
 - Modify the processor to prevent threads from changing its domains (low, high, permission fields)
 - Add a bit to the processor to indicate kernel mode vs. user mode
 - Only kernel mode can change domain register
 - Extend set of permissions for domain to KERNEL-ONLY, and make KERNEL-ONLY permission levels return an exception when in user mode
 - Use handler to process interrupts/exceptions while switching to kernel mode
- Gates and Changing Modes
 - Threads cannot invoke procedures to change between modes in a controlled manner
 - Gates: certain addresses where a thread can enter kernel domain
 - Supervisor call instruction (SVC)
 - Changes processor mode from user to kernel

- Sets an address (into PC) for entry point of the gate manager (to facilitate mode change)
- Gate names are generally numbers; gate manager records gate numbers with their procedures
- SVC instructions must happen all without interruption in order to change modes successfully
 - Processors don't have to perform these as before-or-after actions
- SVC can then return to user mode
 - Changes mode from kernel to user
 - Load program counter from top of stack into PC
- Entering vs. leaving kernel mode
 - Leaving: value loaded in the program counter isn't predefined; kernel sets to saved address
- Special gates also exist to handle interrupts and exceptions