

CS111 Chapter 5.5 Outline #5

- Virtualizing Processors Using Threads
 - Each thread does NOT necessarily have its own processor
 - Thread manager allows sharing of processors
 - One slow thread might slow down the progress of other threads on same processor
- Sharing procedure among multiple threads
 - Thread: abstraction that encapsulates the state of a running module
 - `thread_id ← ALLOCATE_THREAD(starting_procedure, address_space_id)`
 - Not enough space results in an error
 - Allows many threads to share a limited amount of processors
 - Solution depends on the fact that most threads spend most of their lives waiting for events
 - Virtualizing the processor
 - When thread is waiting, processor switches to another thread by saving the previous state and loading the new thread state
 - Also known as time-sharing, processor multiplexing, etc.
 - YIELD: entry point to thread manager, allows sender to add new item (thread?) to buffer
 - Thread manager hands calling thread's processor to some other thread
 - RUNNING (executing on processor) vs. RUNNABLE (ready to run, waiting for processor)
 - Three steps
 - Save: save thread state to resume later
 - Schedule: schedule other thread to run on current processor
 - Dispatch: dispatch processor to thread
 - YIELD procedure:
 - Thread running in thread layer calls YIELD
 - YIELD enters processor layer, which saves its current state
 - When processor later exits processor layer, it runs new thread
 - New thread might run in the same or different address space
 - YIELD is usually written in low-level instructions as a kernel procedure via supervisor call
 - Interrupts invoke interrupt handler, which always runs in processor layer
 - Does not invoke procedures like YIELD
 - No relation to current thread
 - Exceptions occur in thread layer, accessing interrupted thread's state
 - Invokes procedures
 - Specific to current thread
- Implementing YIELD
 - Relies on `GET_THREAD_ID`, `ENTER_PROCESSOR_LAYER`, `EXIT_PROCESSOR_LAYER`, `SCHEDULER`
 - Involves `processor_table` (info on processor) and `thread_table` (array with one entry per thread)
 - Steps
 - Virtualize register `CPUID` to create virtual ID register for each thread
 - YIELD calls `ENTER_PROCESSOR_LAYER`, releasing predecessor
 - Stack pointer is stored
 - Scheduling searches until it finds a runnable thread
 - `EXIT_PROCESSOR_LAYER` loads saved stack pointer from `ENTER`
 - Flow of control results in abandoning the two stack frames allocated to `SCHEDULER` and `EXIT`
 - Thread releases processor by `ENTER` call, then resumes immediately after the same `ENTER` call
 - Known as a "co-routine"
 - Enforcing atomicity: `ENTER` thread acquires `thread_table_lock`; `EXIT` thread releases it
- Creating and terminating threads
 - Supporting management of a variable number of threads requires
 - `EXIT_THREAD()`: destroys and cleans up calling thread to release its state
 - `DESTROY_THREAD(id)`: destroy specified id thread; a thread might need to terminate another
 - At least as many threads as processors need to be supported, so a processor has a separate thread
 - Known as processor-layer thread, or processor thread

- Runs SCHEDULER procedure
 - Each processor needs its own processor thread to deallocate thread layer thread stack spaces
- ALLOCATE_THREAD
 - Allocate memory space for new thread
 - Place empty frame on new stack with EXIT_THREAD as return address
 - Place second empty frame with starting_procedure return address
 - Find FREE entry in thread table
 - Incompletion of these steps returns an error
- DESTROY_THREAD
 - Calling thread cannot free target stack; only processor running that thread can do that
 - Set kill_or_continue variable to KILL
 - Processor thread checks variable during YIELD, possibly releasing resources
- Enforcing modularity with threads: preemptive scheduling
 - Non-preemptive scheduling: thread runs until it gives up processor
 - New thread that called YIELD is then run
 - Problematic: that thread might never give up control
 - Cooperative scheduling (cooperative multitasking)
 - Every thread will call YIELD occasionally
 - Not robust; programmer might forget to put in YIELD, or program errors might arise
 - Preemptive scheduling
 - Thread manager can force a thread to give up the processor based on time (ex: 100 ms)
 - Clock might trigger an interrupt to switch processor layer to kernel mode
 - Interrupt handler invokes exception handler, forcing current thread to YIELD
 - Thread manager needs to allow interrupt handler to invoke procedures
 - But concurrent execution within processor layer with calling YIELD can cause deadlock
 - Thread layer has its own lock mechanism
 - One solution: enable/disable interrupts
 - Make both streams separate (before-or-after actions)
 - Preemptive scheduling isolates thread behavior, making it like each thread has its own process
- Enforcing modularity with threads and address spaces
 - Sharing a single address space between threads can cause accidental memory changes
 - Thread manager needs to switch address spaces when threads are switched
 - One solution: map instructions and data of thread manager into same virtual addresses in every virtual address space
 - Another solution: use hardware to load PMAR, SP, PC as single before-or-after action
 - Returns control to thread in new virtual address space at saved location with saved pointer
- Layering threads
 - Create several threads in thread layer from a single thread in processor layer
 - Interrupt support
 - Processor as a hard-wired thread manager with a processor thread (for SCHEDULER) and interrupt thread (runs interrupt handlers in kernel mode, capable of YIELD)
 - OS layer uses processor threads of processor layer to make a second layer of threads
 - Each application module is given at least one preemptively scheduled virtual processor
 - Each application module can implement its own third-layer thread manager
 - Usually nonpreemptive
 - Threads don't need to be protected, as they belong in the same application module
 - Hardware processor at lowest layer creates a processor thread and an interrupt thread