

CS111 Chapter 5.2.5, 5.6 Outline #9

Chapter 5, Section 5.2.5 (pages 221-222), Chapter 5, Section 5.6 (pages 273-284)

- **Deadlock**
 - Situations just as bad as race conditions
 - Interaction among a group of threads where each thread is waiting for other threads to make progress
 - Wait-for graph can represent deadlocks
 - Nodes are threads/resources (ex: locks)
 - Directed edges go from lock node to thread node (when thread acquires a lock)
 - IFF cycle exists, then threads are deadlocked
 - Preventing deadlocks
 - Enumerate lock usages to ensure all program threads acquire locks in same order
 - Deadlock might also occur when a sender forgets to release/acquire a single lock
 - If buffer is full, receiver will never remove a message from buffer as it can't acquire a lock from the sender, who might in turn be waiting for the receiver to change p.out
 - Avoid simple programming errors
 - Similarly, livelock occurs when a thread repeatedly performs operations but never finishes the sequence
- **Thread Primitives for Sequence Coordination**
 - Thread manager allows threads to share processors
 - Threads can also release processors for other threads to run (much like using sender/receiver YIELD)
 - Polling: thread continually tests a shared variable (in, out) when sender/receiver is scheduled again
 - Undesirable; if test fails, the processor has been acquired/released unnecessarily (unproductive)
 - Get thread manager to schedule only threads with useful work to perform? Sequence coordination
- **Lost Notification Problem**
 - Consider incorrect sender/receiver implementation
 - Includes a shared variable and WAIT and NOTIFY (primitives taking in shared variable)
 - Race condition occurs; notification might be lost when receiver hasn't called WAIT, but when receiver calls WAIT and actually waits, nothing comes
 - Sender continues adding to buffer then WAITs; now both sender/receiver are waiting
 - Operations for testing buffer room, sleep until room, and release shared lock must be before-or-after
 - Can't just lock around operations: who actually owns the lock to create before-or-after actions?
 - Kernel correctness can't depend on user program to set/release lock, nor on user lock's correct implementation
 - Invariant is needed to act as a relation between application-owned and system-owned state
- **Avoiding lost notification problem with eventcounts and sequencers**
 - Designers have suggested using new thread state to characterize event for which thread is waiting under protection of table lock; use event_name to avoid lost notifications
 - Condition variable: holds additional thread state
 - Here we use eventcounts and sequencers, variables to hold additional thread states
 - AWAIT, ADVANCE, TICKET, READ (before-or-after actions)
 - Since AWAIT and ADVANCE are before-or-after, problem doesn't occur
 - Implementing these requires a spinning lock to prevent ACQUIRE from calling AWAIT for an unlocked thread (thread manager is not designed for recursion – no base case)
- **Polling, Interrupts, and Sequence Coordination**
 - Certain threads need to interact with external devices (ex: keyboard manager → keyboard controller)
 - Flip-flop between controller and manager allows activities to be coordinated
 - Keyboard controller sets flip-flop; keyboard manager reads and tests it
 - If it's not set/reads 0, manager yields; if it's set, falls out of loop
 - After reading the flip-flop, it resets to 0 as a side-effect, providing a coordination lock
 - Keyboard manager thread also performs polling, continually checking scheduler for things to run
 - Disadvantages of polling
 - Difficult to predict time until event, so how long should thread poll?
 - Frequent polling thread wastes processor cycles due to lack of user input
 - Devices might require processor to execute managers by a deadline

- Ex: keyboard controller may have a single keystroke to communicate with manager
 - A second keystroke done before the first is absorbed might lose the first
- Similar disadvantages to not using explicit primitives (like AWAIT, ADVANCE)
 - Thread scheduler doesn't know when to run receiver, which might call extra YIELDS
 - When controller has input to be processed, scheduler should be alerted for keyboard manager
- Instead, move polling loop into hardware by using interrupts
 - Keyboard manager enables interrupts (sets processor's interrupt control register to ON) so that processor must take interrupts from controller
 - When processor sees change in flip-flop, it executes interrupt handler
 - Multiple shared flip-flops and an associated map for different handlers allows many interrupts
- Interrupt handler gets processor control within one instruction, so it can meet deadlines
 - Handler could copy keystrokes to a buffer belonging to manager immediately rather than waiting for the chance of it running
- Interrupt handler must carefully read/write shared variables, as it might be called between instructions
 - It can't be share of the thread state currently running on processor
- Interrupt handler must be carefully programmed, since it isn't managed by OS
 - Might cause deadlocks
- Exception handlers are less severe, since handler runs based on the current running thread