

CS111 Chapter 11 Overview, Section 11.2 Outline #11

Chapter 11.2 (11-28 – 11-36), 11.4 (11-49 – 11-54), 11.6 (11-72 – 11-80)

- Authenticating Principles
 - Guard asks questions
 - Who is the principal making the request? Is it actually from that person?
 - Separating trust from authenticating principles
 - Authenticating principals
 - Check real-world identity of principal
 - Physical property (voice recognition); something one has (magnetic cards), knows (PIN)
 - Use verification of identity
 - Cryptographic hash function can protect passwords
 - Secure Hash Algorithm (SHA)
 - Computationally secure (infeasible to break into)
 - Work factor: quantifier used to measure computational security
 - Design principle: minimize secrets
 - Dictionary attack used by brute-forcers (effective, since most passwords suck)
 - Protocol between user and service
 - Authenticates principal to guard
 - Authenticates service to principal
 - Password never travels the network
 - Password only used once per session
- Message Confidentiality
 - Encryption: plaintext → ciphertext
 - Decryption: ciphertext → plaintext
 - ENCRYPT, DECRYPT are implemented using cryptographic transformations
 - Observers cannot translate the messages, even if intercepted
 - Public key and private key
 - Bob hands over public key to Alice, who uses it to encode her message
 - Only Bob has his own private key to decode a message encoded via public key
 - Must be able to withstand attacks like...
 - Ciphertext-only attack: adversary has ciphertext and encrypt/decrypt algorithms; he should not be able to use redundancies and patterns to reconstruct the plaintext
 - Known-plaintext attack: adversary has ciphertext and some of the plaintext; he should not be able to determine the key
 - Chosen-plaintext attack: adversary has ciphertext and chosen corresponding plaintext; he may convince someone to send an encrypted message of plaintext to his liking (adaptive attack)
 - Oracles, which encrypts inputs without question, might encourage adversary attacks
 - Chosen-ciphertext attack: adversary has ciphertext and observes plaintext when it's converted
 - Confidentiality and Authentication
 - Confidentiality: Alice encrypts message
 - Unlikely, confidential info should have a guarantee of where it's from
 - Authentication: Alice signs message
 - Common, since a lot of data is public, but its origin is needed or helpful
 - For both, we must encrypt and sign the message (in this order, to cover up weaknesses)
 - Use different keys for authentication/encryption
 - Explicitness principle: specify who the message is from/for
- Authorization: Controlled Sharing
 - Separate programs that share sensitive data?
 - This restriction on sharing is impossible
 - Rather, use trusted and untrusted networks
 - Still not very convenient
 - Authorization Operations
 - Authorization: gains principal permission to operate on object

- Mediation: operation checks whether principal has operation permission on object
- Revocation: decision that removes previous permission from principal
- Simple Guard Model
 - Authorization matrix with principal rows and object columns
 - Discretionary access-control systems: object creator assigns permissions for that object
 - Non-discretionary: another chosen authority has that responsibility
 - Two instances
 - List systems: organized by column
 - Revocation is less disruptive than ticket system
 - Guard holds list of tokens corresponding to authorized principals
 - Tokens list (access-control list) dictates who has access
 - Requires searching a list, but revocation is easy
 - Ticket systems: organized by row
 - Guard holds ticket (capability) for its object, which it can give out/take back
 - Tickets can get passed around (good and bad), but revocation is hard
 - Tokens and tickets must be shielded from forgery
 - Agency allows a switch between list and ticket systems
 - Protection groups: principals used by more than one user
 - Might be implemented using an ACL
 - Kind of like a permission group
 - UNIX principals identified by strings or integer names (UIDs)
 - GID is the principal for whom the process is running
 - Use principal of least privilege
 - Authentication users based on UID, GID, and password
 - Access Control Check
 - Superuser (UID 0) has permissions by default
 - Matching UID of owner and process prompts for ACL check
 - Matching GID of process and file prompts for ACL check
 - UID/GID don't match, so we check ACL for other matches
- Caretaker Model
- Flow-Control Model