

CS111 Two Articles Outline #14**An Introduction to Device Drivers; Understanding Modern Device Drivers**

- Device Drivers: entry point for kernel hackers to learn the code without being overwhelmed by complexity
 - Modularity: black box that hides details; enables hardware to respond to a well-defined interface
 - Easy to write drivers
 - Drivers can be built separately from the kernel and “plugged in”
 - Turnover rate of hardware technology is fast; adding a driver means more users will use a hardware
 - Also, the source can be released
- Role of Device Driver
 - Provides mechanism, not policy (and is therefore flexible)
 - Mechanism vs policy: capabilities provided vs. how capabilities are used
 - Graphic display example: server vs. window/session managers
 - ftpd: file transfer mechanism
 - Better to have policy-free drivers, since they can be used more freely for different purposes
 - Applications themselves can implement policy (floppy vs. floppy permissions/properties, etc.)
 - Easier to implement and easier to use
 - Flexible driver needs to have capabilities without constraints
 - Sometimes policy is necessary so the hardware is used correctly (ex: handle certain bits in I/O driver)
 - Device driver sits between hardware and application; programmer has control over device appearance
 - Tradeoff between providing options to user versus keeping it simple
 - Programmer has freedom in implementing concurrency, memory mapping, etc.
- Splitting the Kernel
 - Concurrent processes do different tasks
 - Kernel (executable code that handles resource requests) is split into different roles
 - Process management
 - Creates/destroys processes, connects them outside, schedules CPU sharing in processes
 - Memory management
 - Policies for dealing with memory, builds virtual address space
 - Filesystems
 - Structured filesystem on top of hardware; types include ext3, FAT, etc.
 - Device control
 - Coordinates device drivers for operations mapped to a physical device
 - Networking
 - Packets are collected, identified, and dispatched; controls programs’ network aspect
- Loadable Modules
 - Ability to extend/remove kernel features at runtime
 - This code (modules) can be linked/unlinked to kernel by insmod/rmmod
- Classes of Devices and Modules
 - Better use a module per functionality to allow decomposition for scalability/extendibility
 - Character devices
 - Implements behavior concerning streams of bytes (open, close, write, read)
 - Block devices
 - Device that hosts a filesystem; transfers many bytes at once
 - Network interfaces
 - Interface facilitates network transactions to exchange data between hosts
 - Universal Serial Bus (USB)
 - Filesystems can also determine the amount of info in a block device, etc.
 - Software driver that maps low-level and high-level data structures
- Security Issues
 - Security checks in system are enforced by kernel code
 - Security policy best handled under kernel at high level
 - Driver writers must know when to put security policy that prevents adverse effects
 - They must also avoid bugs (buffer overflow, etc.)

- Be careful when running precompiled binaries
- Version Numbering
 - Every Linux software has its own release number
 - Even number kernels are stable for general distribution; odd ones are development and ephemeral
- License Terms
 - Linux is under version 2 of GNL General Public License (GPL)
 - Allows growth of knowledge by allowing everyone to modify program
- Understanding Modern Device Drivers
 - Device drivers are biggest OS contributors
 - Driver code:
 - Common assumptions: characteristics of driver code functionality?
 - Interactions: How do they interact with kernel, devices, and buses?
 - Contents: Can we reduce complexity/driver size using libraries?
- Introduction
 - Device drivers are 70% of Linux code base
 - Understanding driver code to improve reliability and reduce complexity
- Background
 - Communication request and OS service access: driver and kernel
 - Executing operations: driver and device
 - Communication management: driver and bus
 - Driver/Device
 - Character drivers (byte-stream), block drivers (random block access), network drivers (packet streams)
 - Driver Research Assumptions
 - Interaction assumptions refer to how driver interacts with kernel
 - Architecture assumptions refer to role of the driver
 - Some assumptions
 - Drivers support single chipset (like efforts to synthesize drivers from formal specs)
 - Drivers are conduit of communicating data and signaling device; it does little processing
- What do Drivers Do?
 - Translates between high level inputs to low level, hardware instructions (I/O)
 - Not fully accurate...
 - Methodology
 - DrMiner (static analysis tool) uses tagging to label entry points of driver code
 - Function breakdown of driver code
 - Mostly initialization and cleanup; only 23% of driver code request handles and interrupts
 - Optimization therefore depends on initialization and configuration
 - Research needs to look more into media, GPU, and wireless drivers (rapidly changing ones)
 - Usually we focus on generally important ones like Ethernet and sound
 - Do drivers belong to classes?
 - Most driver functionality falls into class behavior, but quite a few do not
 - Do drivers do significant processing?
 - Yes, a substantial amount do some processing
 - How many device chipsets does a single driver support?
 - Multiple chipsets per driver; efficient; avoid system that uses unique drivers for chipsets
 - This increases complexity, expands driver code
 - Most assumptions are correct, but not all the time
- Driver Interactions
 - Four topics
 - Kernel resources consumed by driver
 - How/when drivers interact with devices
 - Differences in driver structure across I/O buses
 - Threading/synchronization model used by driver code

- Driver/kernel interaction
 - Classified into one of five
 - Kernel library (generic support routines, standard data structures)
 - Memory management (allocation)
 - Synchronization (locks)
 - Device library (subsystem libraries supporting class of device or I/O)
 - Kernel services (access to other subsystems like files, memory, scheduling)
 - A large amount is generic routines
- Driver/device interaction
 - Perform I/O is a function that does I/O on itself or calls I/O function
 - Number and type of device interactions vary, so their cost varies as well
- Driver/bus interaction
 - Drivers are meant for devices that attach to some PCI (or USB)
 - PCI devices' flexibility/performance come at cost of complexity and less standardized interface
- Driver Concurrency
 - Drivers are required to multiplex access to device
 - Threaded code needs to be run in kernel (hard to do otherwise)
- Driver Redundancy
 - Related devices share code with a small amount of per-device code
 - Reducing complexity
 - Reduce driver shape down to single signature value
 - There are many opportunities to reduce volume of code by abstraction
 - Procedural abstractions for driver sub-classes
 - Better multiple chipset support
 - Table driven programming
 - Procedural abstraction is most useful (moves shared code to library and provides parameters)