

CS111 Sections of Chapter 4, Article Outline #18

Chapter 4, section 4.5, pages 184-195. [What Cloud Computing Really Means](#), Eric Knorr, Galen Gruman, Infoworld.

- Case Study: Network File System (NFS)
 - Network File System is client/service application that gives shared file storage for clients on the network
 - NFS client puts remote file system onto client's local system much like a local UNIX file system
 - Purpose was to share files for collaboration
 - Simplifies management of workstation collection
 - Admin doesn't have to manage or make backups for each workstation individually
 - Also saves the cost of buying a disk (interface) for every workstation
 - Design goals for NFS
 - Work with existing applications (use same local UNIX semantics)
 - Easily deployable and easy to retrofit into existing UNIX systems
 - Client implementable in other operating systems (can't just be UNIX, but also allow Windows)
 - Efficient enough for users to tolerate, but doesn't need as high performance as in local systems
- Naming Remote Files and Directories
 - Appears the same way as a local UNIX file system
 - A "mounter" program mounts (or grafts) remote file system onto local name space
 - Mounter sends RPC to file server host to obtain a "file handle" (32-byte object name)
 - File handle IDs object on remote NFS server; invisible to applications, but internal to name files
 - File handle is structured: contains file system identifier, inode number, and generation number
 - Generation number: server uses to locate file
 - File system identifier: server IDs file system responsible for file
 - inode number: lets the IDed file system locate file on disk
 - Path names are not included in file handles
 - This is to keep NFS and local file systems consistent in opening files (in case some are renamed)
 - To reuse inode numbers, we increase generation number
 - In ambiguous situations, NFS gives "stale file handle" error for simpler implementation
 - File handles are available even across server failures
- NFS Remote Procedure Calls
 - Opening a remote file "f" involves NFS client sending RPC: LOOKUP(dirfh, "f")
 - NFS server uses extracted identified file system and inode number to locate directory's inode
 - NFS then uses directory's inode to find "f" and creates its file handle for the client
 - Closing a remote file "f" that has not been modified doesn't require RPC
 - NFS RPC hold their own information in order to be stateless (server maintains only on-disk file states)
 - Statelessness makes server failure recovery simple; client just repeats the request (idemponent)
 - Sometimes multiple RPC calls might give inconsistent results if a network reply is lost
 - Example: NFS client sends REMOVE RPC; server removes file, loses reply; NFS REMOVES again, but there's already no file
 - NFS should also avoid issuing multiple RPCs when there's no server failure
 - Use a "soft" state maintained by server that goes away if server fails (AKA a reply cache)
 - Reply cache maintained as soft state, since non-volatile storage is expensive
 - When given a request, server looks up transaction ID in reply cache
 - If it's there, return cache reply; otherwise, process the request
 - UNIX maintains state, but NFS is stateless, as it tries to simplify recovery
 - This is one instance where NFS has different behavior
- Extending the UNIX File System to Support NFS
 - To get NFS as UNIX extension, designers use vnodes (virtual nodes) to split file system
 - vnode is volatile memory structure that abstracts whether file/directory is local or remote
 - File system caller doesn't have to worry about distinguishing local/remote in files or directories
 - Operating on a file involves calling the procedure through vnode interface
 - To open a remote file, caller invokes PATHNODE_TO_VNODE
 - Passes vnode for current working directory and path name for the file

- LOOKUP is invoked in the vnode layer for each path name component; local-implementation LOOKUP is called for local files and a remote-implementation LOOKUP is called for remote files
 - NFS server extracts file system identifier and inode from file handle; LOOKUP is invoked in the vnode layer and a vnode (if local file is present) is returned to NFS
 - NFS client then creates a vnode with file handle on client computer, returning it to file system call layer on machine
 - On file system call resolution, file descriptor for the file is returned
- Caches are utilized to store vnode for open files and recently-used vnodes (and their attributes)
- Caching reduces latency by reducing RPC cost for a client
 - Clients then use fewer RPCs, so a server can support more clients
 - If clients cache same file, we need to ensure R/W coherence
- Coherence
 - NFS could guarantee R/W coherence for either all operations or just some
 - Close-to-open consistency: only provide R/W coherence on OPEN, WRITE, and CLOSE
 - After OPEN/WRITE/CLOSE, a second client using the same file observes the first WRITE
 - Or every READ/WRITE provides coherence; every WRITE results in a READ that observes results
 - Allows higher data rates for reading/writing big files
 - Client can send several reads without waiting for response
 - In the second open-to-close consistency implementation above, no guarantees
 - READ may return data before or after last WRITE when same file is open twice
 - Client WRITES by modifying local cached version, so no overhead of RPCs
 - CLOSE then sends cached writes to server
 - Delaying modified blocks until CLOSE absorbs overwritten modifications
 - This R/W coherence allows user to edit program on one computer and compile it on a different one
 - After file is modified/written, compiler observes all the edits
 - However, multiclient DB programs that read or write records stored in NFS file might not work correctly due to concurrent execute operations
 - Also, blocks are cached instead of entire files; fetching a file may result in intermixed blocks
- NFS Version 3+ supports 64-bit numbers, adds asynchronous write, etc.
- Version 4 maintains some state, prevents intruders from snooping/modifying network traffic, has more efficient close-to-open consistency scheme, and works well over internet
- What cloud computing really means
 - Cloud computing: virtual servers available over internet (an updated version of utility computing)
 - Adds capacity or capabilities on the fly without new infrastructure, personnel, or licenses
 - May be subscription-based or pay-per-use service; for example, storage services to spam filtering
 - SaaS: Software as a Service
 - Browser delivers an application via multitenant architecture (HR apps, even “desktop” Google Apps)
 - Customer has no investment in services or licenses; providers just have one cheap app to maintain
 - Utility computing
 - Offer solutions like storage and virtual services; often used by IT on demand
 - Used on Amazon, Sun, IBM, etc.
 - Web services in the cloud
 - Like SaaS, these offer APIs for developers to use functionality over internet
 - Range from business services to Google Maps API
 - Platform as a service
 - Another SaaS variation that gives development environment as service (Google App Engine)
 - MSP (managed service providers)
 - IT (not end-user) application like virus scanning, email, or application monitoring service (very old)
 - Service commerce platforms
 - SaaS and MSP hybrid; service hub for user interaction, like a trading environment (travel services)
 - Internet integration
 - Integration of cloud-based services (serving SaaS providers, using in-cloud integration techniques)
- More accurately “sky computing” due to isolated clouds of services which IT customers “plug into individually”
- Loosely coupled services operating on an infrastructure that is agile and scalable